

# RPC-based Proxy Server Web Cache

The end goal of this project is to write the components of a proxy-server web cache and to experiment with various cache policies. Along the way, you will learn curl's C interface and program remote procedure calls (RPC) with the rpcgen tool.

## Contents

- 1 [Setup](#)
- 2 [Collaboration](#)
- 3 [RPC Tasks](#)
  - 3.1 [RPC Client](#)
  - 3.2 [RPC Server and libcurl](#)
- 4 [Cache Policy Implementations](#)
- 5 [Cache Policy Experiments](#)
- 6 [Grading](#)

## Setup

1. Gain access to a well-maintained linux machine. If you do not have this already, then it is recommended that you follow the [instructions](#) for downloading a installing VirtualBox and the AOS Virtual Image.

2. Download the proxy-server project starter code [proxy-server.tar.gz](#)

3. Check that that rpcbind daemon is running on your system with

```
ps -e | grep rpcbind
```

4. Check that libcurl is installed with

```
which curl-config
```

## Collaboration

For this and all projects, students must work on their own.

## RPC Tasks

### RPC Client

Your first task is to complete a toy program using rpcgen. Start by defining the desired interface in the file proxy\_rpc.x. For this assignment, we need a single procedure called httpget\_1. The file is written in a special RPC language and is compiled into C code that handles the client-server connection. Consult the [rpcge Programming Guide](#) for details. Running

```
rpcgen proxy_rpc.x
```

generates three files

- proxy\_rpc\_clnt.c — the code for the client side
- proxy\_rpc\_svr.c — the code for the server side
- proxy\_rpc.h — a header file useful to both None of files should be modified. They are intended to be linked with other files, as we will do in the remainder o exercise.

Next, write a client program that makes a remote procedure call to httpget\_1 and prints the string that is returned. Detailed instructions can be found in [client.c](#) / you will find the [rpcgen Programming Guide](#) useful. Compile your code with

```
rpcgen proxy_rpc.x
make client_test client
```

Then run

```
sudo ./client_test
```

in one terminal window (depending on how rpcbind is setup you may or may not need sudo), and

```
./client localhost example.com
```

in another. The window where you ran the client program should print “SUCCESS!”

#### Files for submission

- proxy\_rpc.x
- client.c

#### Support files:

- Makefile
- client\_test.c

---

## RPC Server and libcurl

Now we consider the proxy server. Inside the file proxy.c you need to define the httpget\_1 procedure with the correct signature. Again, the rpcgen [Programming will be useful](#).

To actually get the contents of the requested url, you will use libcurl. The [libcurl project site](#) contains more information and more examples. A particularly useful example will be [getinmemory.c](#). The Makefile already links proxy.c to the libcurl library. Replacing client\_test with proxy in the instructions for the rpc client above print of the data returned from the provided url in the client’s window.

#### Files for submission:

- proxy.c

#### Support files:

- Makefile

---

## Cache Policy Implementations

Next, you will implement four different replacement policies for the webcache.

- Least Recently Used (LRU) — remove the item that has been requested least recently.
- Random — remove an item chosen uniformly at random from the cache.
- Least Frequently Used (LFU) — remove the item that has been used the fewest times from the cache.
- LRUMIN — Group the cache entries by size on a base-2 logarithmic scale and remove entries from the group closest to the size needed. Precise specification can be found in the comments of [lrumin.c](#)

Write your code for these policies in the files lru.c, rnd.c, lfu.c, and lrumin.c, respectively, all of which implement the API given in [gtcache.h](#). All implementations must be efficient. Inefficiency will result in points lost.

To help you implement these policies, the following data structures are provided.

- hstbl.[ch] — contains a string-indexed hashtable to help you quickly retrieve data from the cache.
- indexminpq.[ch] — contains an indexed minimum priority queue[link] generic data structure to help you quickly figure out which item to replace in the cache.
- indexrndq.[ch] — contains an indexed randomized queue supporting constant-time arbitrary deletion.
- steque.[ch] — contains a steque data structure which might be useful for keeping track of available ids.

If you need to record the time a request is made, use the gettimeofday method in sys/time.h. While testing your code, you do not want to actually use curl and make a web request. The file webdata.txt contains a list of 95 urls and the content lengths of the responses and the files fake\_www.[ch] contain some utilities for simulating web requests. You may find these helpful as you write your testcode.

#### Files for submission:

- lru.c
- rnd.c
- lfu.c
- lrumin.c

**Support files:**

- Makefile
- gtcache.h
- hshtable.[ch]
- indexminpq.[ch]
- indexrndq.[ch]
- steque.[ch]

## Cache Policy Experiments

No one cache policy is best for all request patterns. To illustrate this point, you are to provide lists of requests that affect the hit rates of the cache algorithms in particular ways.

The program `cache_perform.c` will be used to evaluate your data over cache sizes between 1 and 16 MB. Compile your code with

```
make perform
```

which will produce `lru_perform`, `rnd_perform`, `lfu_perform`, and `lrumin_perform` executables. If you list of requests is in a file `requests.txt`, you can then you can test the performance of your algorithms by passing the file through `stdin`. For example, to see the performance of `lru` on `requests.txt` with 2000 requests, one would run

```
./lru_perform webdata.txt 2000 < requests.txt
```

The `requests.txt` file should contain one url per line. When the end of the list is reached `cache_perform` starts over until 2000 requests have been made.

We will say that one policy “beats” another if it outperforms it on a majority of the cache sizes between 1 and 16MB that we test. For simplicity of grading, your data will be run against our implementations of these policies. You are to submit four request pattern files. The requirements are indicated by the filename.

**Files for submission:**

- `rnd_beats_lru.txt`
- `lfu_beats_lru.txt`
- `lru_beats_lfu.txt`
- `lrumin_beats_lru.txt`

**Support files:**

- Makefile
- `cache_perform.c`
- `gtcache.h`

## Grading

Deliverable	Percentage
<code>client.c</code> , <code>proxy_rpc.x</code>	15%
<code>proxy.c</code>	15%
<code>lru.c</code>	10%
<code>rnd.c</code>	8%
<code>lfu.c</code>	10%
<code>lrumin.c</code>	10%
<code>rnd_beats_lru.txt</code>	8%
<code>lfu_beats_lru.txt</code>	8%
<code>lru_beats_lfu.txt</code>	8%
<code>lrumin_beats_lru.txt</code>	8%

INFORMATION

[Help and FAQ](#)

COMMUNITY

[Blog](#)

[Meetups](#)

[News & Media](#)

[Developer API](#)

UDACITY

[About](#)

[Jobs](#)

[Legal](#)

FOLLOW US ON

© 2011-2014 Udacity, Inc.